

DIFFERENCES BETWEEN HIVE AND RDBMS

HIVE VS RDBMS

HIVE IS A DATA WAREHOUSE ON TOP
OF HADOOP

HIVE VS RDBMS

**DATAWAREHOUSES ARE INHERENTLY
DIFFERENT FROM RELATIONAL
DATABASES**

HIVE VS RDBMS

THEY ARE GEARED TOWARDS
ANALYTICAL PROCESSING

HIVE VS RDBMS

RDBMS ARE GEARED TO
TRANSACTION PROCESSING

HIVE VS RDBMS

WHILE THERE ARE MANY
DIFFERENCES IN THESE 2 PARADIGMS

LET'S CONCENTRATE ON 2 SPECIFIC
DIFFERENCES

HIVE VS RDBMS

2 SPECIFIC DIFFERENCES

BATCH PROCESSING

SCHEMA-ON-READ

HIVE VS RDBMS

2 SPECIFIC DIFFERENCES

BATCH PROCESSING

SCHEMA-ON-READ

HIVE VS RDBMS

BATCH PROCESSING

HIVE IS OPTIMAL FOR PROCESSING
REALLY LARGE-SCALE DATASETS

GIGABYTES/PETABYTES

HIVE VS RDBMS

BATCH PROCESSING

THE DATA IN HIVE IS MEANT TO BE
USED FOR ANALYTICAL PURPOSES

HIVE VS RDBMS

BATCH PROCESSING

FOR EX: TO PROCESS ORDERS DATA
FOR TRENDS IN LAST 3 YEARS

HIVE VS RDBMS

BATCH PROCESSING

SINCE HIVE USES HADOOP, MAPREDUCE

IT LEVERAGES THE TREMENDOUS PARALLEL
COMPUTING POWER OF THESE FRAMEWORKS

HIVE VS RDBMS

BATCH PROCESSING

TRADITIONAL DATABASES IMPROVE
PERFORMANCE USING INDEXES

HIVE VS RDBMS

BATCH PROCESSING

BUILDING INDEXES TAKE A LOT OF
TIME AND OCCUPIES DISK SPACE

HIVE VS RDBMS

BATCH PROCESSING

INCREASING DISK SPACE ON A SINGLE
SYSTEM IS AN EXPENSIVE AFFAIR

HIVE VS RDBMS

BATCH PROCESSING

IN ADDITION PERFORMANCE DOES NOT
INCREASE LINEARLY WITH DISK SPACE

HIVE VS RDBMS

BATCH PROCESSING

ON THE OTHER HAND, HIVE/HADOOP USE LARGE
NUMBERS OF CHEAP MACHINES IN PARALLEL

HIVE VS RDBMS

BATCH PROCESSING

DOUBLING PERFORMANCE SIMPLY
MEANS YOU NEED TO DOUBLE THE
NUMBER OF MACHINES

HIVE VS RDBMS

BATCH PROCESSING

HIVE DOES NOT PERFORM WELL FOR
TRANSACTIONAL PROCESSING

HIVE VS RDBMS

BATCH PROCESSING

EVEN FETCHING A SINGLE ROW WILL
LAUNCH A MAP-REDUCE JOB THAT
MIGHT TAKE MINUTES TO RUN

HIVE VS RDBMS

BATCH PROCESSING

A WELL DESIGNED RDBMS CAN
ANSWER TRANSACTIONAL QUERIES IN
MILLI/MICROSECONDS

HIVE VS RDBMS

BATCH PROCESSING

HIVE ALSO DOES NOT ALLOW ROW
LEVEL UPDATES IN ITS TABLES

HIVE VS RDBMS

BATCH PROCESSING

ONCE THE DATA IS WRITTEN TO HIVE,
ITS PURPOSE IS READ-ONLY

HIVE VS RDBMS

BATCH PROCESSING

THIS IS TOTALLY FINE FOR ANALYTICAL PROCESSING

ANALYSTS WOULD ONLY READ THE
HISTORICAL DATA, NEVER UPDATE IT

HIVE VS RDBMS

BATCH PROCESSING

WHY ARE UPDATES NOT ALLOWED IN HIVE?

HIVE IS BASED ON SCHEMA-ON-READ

HIVE VS RDBMS

2 SPECIFIC DIFFERENCES

BATCH PROCESSING

SCHEMA-ON-READ

SCHEMA-ON-READ

THE SCHEMA IS THE **DESCRIPTION** OF
A DATABASE TABLE

COLUMN NAMES,
COLUMN TYPES
CONSTRAINTS

SCHEMA-ON-READ

COLUMN NAMES,
COLUMN TYPES
CONSTRAINTS

THESE 3 ARE
DEFINED WHEN YOU
CREATE A TABLE

SCHEMA-ON-READ

COLUMN NAMES,
COLUMN TYPES
CONSTRAINTS

WHEN YOU READ
FROM OR WRITE TO
A TABLE, YOU EXPECT
THAT THIS SCHEMA IS
ENFORCED

SCHEMA-ON-READ

HOW IS THE SCHEMA ENFORCED?

TRADITIONAL DATABASES DO IT USING

SCHEMA-ON-WRITE

HIVE USES

SCHEMA-ON-READ

TRADITIONAL DATABASES

SCHEMA-ON-WRITE

TRADITIONAL DATABASES HAVE
COMPLETE CONTROL OVER THE
DATA STORAGE

TRADITIONAL DATABASES

SCHEMA-ON-WRITE

NO EXTERNAL PROGRAM CAN
ACCESS THE DATA WITHOUT GOING
THROUGH THE DATABASE

TRADITIONAL DATABASES

SCHEMA-ON-WRITE

THE DATABASE IS THE
SOLE GATEKEEPER

TRADITIONAL DATABASES

SCHEMA-ON-WRITE

THE DATABASE ENFORCES ACID
PROPERTIES

ATOMICITY, CONSISTENCY,
ISOLATION, DURABILITY

TRADITIONAL DATABASES

SCHEMA-ON-WRITE

THIS MEANS THAT THE DATABASE
ENFORCES THE SCHEMA OF A TABLE
WHEN DATA IS WRITTEN

TRADITIONAL DATABASES

SCHEMA-ON-WRITE

NO DATA CAN EXIST IN THE TABLE
WITHOUT FOLLOWING THE SCHEMA

SCHEMA-ON-READ

HOW IS THE SCHEMA ENFORCED?

TRADITIONAL DATABASES DO IT USING

SCHEMA-ON-WRITE

HIVE USES

SCHEMA-ON-READ

HIVE

SCHEMA-ON-READ

IN HIVE THE UNDERLYING DATA OF
TABLES IS STORED AS FILES IN HDFS

HIVE

SCHEMA-ON-READ

**HIVE IS NOT THE SOLE OWNER OF THE
FILES**

HIVE

SCHEMA-ON-READ

**THE FILES CAN BE USED AND MODIFIED
BY OTHER CLIENTS**

HIVE

SCHEMA-ON-READ

**FOR EXAMPLE, YOU CAN HAVE THE
SAME FILES SHARED BETWEEN HIVE,
HBASE AND CASSANDRA**

**THESE ARE DIFFERENT DATABASE
TECHNOLOGIES THAT WORK USING HADOOP**

HIVE

SCHEMA-ON-READ

**BECAUSE THE UNDERLYING FILES CAN
BE CHANGED AT ANY TIME, HIVE
CANNOT ENFORCE**

SCHEMA-ON-WRITE

HIVE

SCHEMA-ON-READ

INSTEAD, HIVE ENFORCES

SCHEMA-ON-READ

HIVE

SCHEMA-ON-READ

IN THE HIVE METASTORE, THERE ARE
INSTRUCTIONS FOR HIVE ON HOW TO
READ AND PARSE THE HDFS FILES

HIVE

SCHEMA-ON-READ

IN THE HIVE METASTORE, THE
SCHEMA OF A TABLE IS ALSO STORED

HIVE

SCHEMA-ON-READ

WHEN YOU READ THE DATA IN HIVE, IT
WILL PARSE THE FILE AND TRY TO
IMPOSE THE SCHEMA

HIVE

SCHEMA-ON-READ

HIVE MAY NOT ALWAYS SUCCEED IN
IMPOSING THE SCHEMA

HIVE

SCHEMA-ON-READ

WHAT IF YOU HAD A FILE WITH LINES

VITTHAL, SRINIVASAN

JANANI, RAVI

SWETHA, HANURAJ

HIVE TRIES TO IMPOSE THE SCHEMA

FIRSTNAME, LASTNAME, AGE

HIVE

SCHEMA-ON-READ

FIRSTNAME, LASTNAME, AGE

VITTHAL, SRINIVASAN

JANANI, RAVI

SWETHA, HANURAJ

THE FILE HAS ONLY 2 COLUMNS

HIVE

SCHEMA-ON-READ

FIRSTNAME, LASTNAME, AGE

VITTHAL, SRINIVASAN

JANANI, RAVI

SWETHA, HANURAJ

**HIVE EXPECTS THE ROW TO HAVE 3
COLUMNS BUT FINDS ONLY 2 COLUMNS**

HIVE

SCHEMA-ON-READ

FIRSTNAME, LASTNAME, AGE

VITTHAL, SRINIVASAN

JANANI, RAVI

SWETHA, HANURAJ

**IF YOU TRY TO QUERY THE AGE COLUMN,
IT WILL RETURN NULL VALUES**

HIVE

SCHEMA-ON-READ

**HIVE DOES IT'S BEST TO IMPOSE THE
SCHEMA ON A FILE**

HIVE

SCHEMA-ON-READ

BECAUSE HIVE IS SCHEMA-ON-READ

**IT DOES NOT CHECK THE DATA AT ALL
WHEN DOING WRITE OPERATIONS**

HIVE

SCHEMA-ON-READ

DURING LOAD/INSERT OPERATIONS, HIVE
WILL JUST DUMP THE DATA INTO A FILE
WITHOUT CHECKING THE SCHEMA

HIVE

SCHEMA-ON-READ

**SINCE IT IS UNAWARE OF SCHEMA
DURING WRITE OPERATIONS, YOU
CANNOT DO ROW LEVEL UPDATES/DELETES**

HIVE

SCHEMA-ON-READ

FOR THE SAME REASON, HIVE CANNOT
SUPPORT CONSTRAINTS LIKE PRIMARY
KEY, FOREIGN KEY ETC

DIFFERENCES BETWEEN HIVEQL AND SQL

CONSTRAINTS

JOINS

SUBQUERIES

FUNCTIONS

DELETE/UPDATE

ALL OF THESE ARE FEATURES IN SQL

CONSTRAINTS

FUNCTIONS

JOINS

DELETE/UPDATE

SUBQUERIES

**HIVEQL ALSO PROVIDES MANY OF
THESE FEATURES...**

CONSTRAINTS

JOINS

.. WITH SOME LIMITATIONS

FUNCTIONS

SUBQUERIES

DELETE/UPDATE

CONSTRAINTS

IN SQL YOU CAN USE CONSTRAINTS TO
MAINTAIN RELATIONSHIPS BETWEEN TABLES

FOREIGN KEY CONSTRAINTS

YOU CAN ALSO IMPOSE
CONSTRAINTS ON DATA IN A COLUMN
PRIMARY KEY, NOT NULL, UNIQUE

CONSTRAINTS

FOREIGN KEY CONSTRAINTS

PRIMARY KEY, NOT NULL, UNIQUE

**HIVE DOES NOT
SUPPORT ANY OF THESE
CONSTRAINTS**

CONSTRAINTS

CONSTRAINTS ARE
ENFORCED BY A DATABASE
WHEN DATA IS WRITTEN

CONSTRAINTS

HIVE DOES NOT CARE
ABOUT THE SCHEMA OF
THE DATA WHEN IT'S
WRITTEN

CONSTRAINTS

HIVE IS SCHEMA-ON-READ, WHICH
MEANS IT ONLY IMPOSES
SCHEMA WHEN IT READS A TABLE

CONSTRAINTS

ONE OTHER PURPOSE OF
CONSTRAINTS TRADITIONALLY IS
TO SPEED UP QUERY EXECUTION

CONSTRAINTS

HIVE HAS LIMITED SUPPORT
FOR INDEXES TO SPEED UP
QUERY EXECUTION

CONSTRAINTS

HIVE PRIMARILY USES FEATURES
LIKE **PARTITIONING AND BUCKETING**
IN IMPROVE QUERY PERFORMANCE

CONSTRAINTS

JOINS

SUBQUERIES

VIEWS

FUNCTIONS

DELETE/UPDATE

IN SQL YOU CAN DELETE/
UPDATE ROWS THAT
SATISFY A CONDITION

DELETE/UPDATE

~~delete from employees where
firstname="Kiran";~~

~~update employees set
lastname="Srinivasan"
where firstname="Vittthal";~~

DELETE/UPDATE

THESE DO NOT WORK IN HIVE

CONSTRAINTS
ONCE AGAIN THIS IS BECAUSE
OF SCHEMA-ON-READ
FUNCTIONS
SUBQUERIES

DELETE/UPDATE

**DELETE AND UPDATE ARE
WRITE OPERATIONS**

DELETE/UPDATE

CONSTRAINTS

**HIVE DOES NOT IMPOSE ANY SCHEMA
DURING WRITE OPERATIONS**

DELETE/UPDATE

CONSTRAINTS

JOINS

SUBQUERIES

FUNCTIONS

DELETE/UPDATE

HIVE AND SQL BOTH
HAVE SEVERAL
FUNCTIONS TO
PROCESS DATA

FUNCTIONS

DELETE/UPDATE

AVG, COUNT, MONTH, YEAR
ETC

HIVE HAS MANY
MORE BUILT-IN
FUNCTIONS THAN SQL
FOR SOME USE CASES

FUNCTIONS

DELETE/UPDATE

MATHEMATICAL FUNCTIONS
LIKE COVARIANCE

**HIVE HAS MANY
MORE BUILT-IN
FUNCTIONS THAN SQL
FOR SOME USE CASES**

FUNCTIONS

DELETE/UPDATE

**FUNCTIONS FOR STATISTICAL
PROCESSING AND DATA MINING**

CONSTRAINTS

JOINS

SUBQUERIES

FUNCTIONS

DELETE/UPDATE

CONSTRAINTS

JOINS

FUNCTIONS

SUBQUERIES

**HIVE HAS EXTENSIVE
SUPPORT FOR JOINS**

DELETE/UPDATE

CONSTRAINTS

JOINS

FUNCTIONS

SUBQUERIES

**IT EVEN ALLOWS THE USER TO
CONTROL THE EXECUTION OF
THE JOIN TO SOME EXTENT**

DELETE/UPDATE

CONSTRAINTS

JOINS

FUNCTIONS
SUBQUERIES
DELETES/UPDATES
**THIS ALLOWS USERS TO
OPTIMIZE JOINS FOR
CERTAIN USE CASES**

CONSTRAINTS

JOINS

**THERE ARE 2 MAIN
DIFFERENCES IN JOINS
BETWEEN HIVE AND SQL**

CONSTRAINTS

JOINS

FUNCTIONS

SUBQUERIES

1. HIVE SUPPORTS EQUIJOINS ONLY

DELETE/UPDATE

2. HIVE DOES NOT SUPPORT
NATURAL JOINS

EQUIJOINS

CONSTRAINTS

JOINS

HERE IS A JOIN QUERY

```
select a.firstname, b.subordinate  
from employees join subordinates on  
a.empId=b.empId;
```

CONSTRAINTS

EQUIJOINS

JOINS

FUNCTIONS

select a.firstname, b.subordinate
from employees join subordinates on
a.empId=b.empId;

DELETE/UPDATE

IN SQL THIS CONDITION
CAN BE ANYTHING

CONSTRAINTS

EQUIJOINS

JOINS

```
select a.firstname, b.subordinate  
from employees join subordinates on  
a.empId=b.empId;
```

IN HIVE THIS CONDITION
HAS TO BE AN EQUALITY

NATURAL JOINS

JOINS

```
select a.firstname, b.subordinate  
from employees, subordinates;
```

**NATURAL JOINS ARE JOINS WHERE
THE JOIN CLAUSE IS IMPLICIT**

NATURAL JOINS

JOINS

```
select a.firstname, b.subordinate  
from employees, subordinates;
```

**THE TABLES ARE JOINED IMPLICITLY
BASED ON COMMON COLUMNS**

CONSTRAINTS

NATURAL JOINS

JOINS

FUNCTIONS

SUBQUERIES

HIVE DOES NOT SUPPORT

VIEWS

NATURAL JOINS

TEMPORAL UPDATE

CONSTRAINTS

JOINS

FUNCTIONS

SUBQUERIES

**HIVE SUPPORTS
SUBQUERIES**

DELETE/UPDATE

CONSTRAINTS

JOINS

FUNCTIONS

SUBQUERIES

DELETE/UPDATE

**SOME SPECIFIC TYPES OF
SUBQUERIES ARE NOT ALLOWED**

CONSTRAINTS

JOINS

SUBQUERIES

UNION IS SUPPORTED
BUT INTERSECT AND
MINUS CONSTRUCTS
DON'T EXIST IN HIVE

CONSTRAINTS

JOINS

FUNCTIONS

SUBQUERIES

THERE ARE SOME RESTRICTIONS
ON HOW IN/EXISTS ARE USED

