

**VIEWS**

# VIEWS

LET'S SAY THERE IS A REALLY COMPLICATED  
QUERY YOU USE AGAIN AND AGAIN

YOU COULD **SAVE THE QUERY AS A VIEW**

# VIEWS

YOU COULD THEN QUERY THAT VIEW  
EXACTLY AS IF IT WERE A TABLE

RATHER THAN HAVE IT AS A COMPLICATED  
SUB-QUERY IN YOUR OTHER QUERIES

# HERE IS A PRETTY COMPLEX QUERY

SELECT PRODUCTNAME,REVENUE FROM (

SELECT

P.PRODUCTNAME, YEAR (ORDERDATE) AS ORDERDATE, SUM (REVENUE) AS REVENUE

FROM

(SELECT \* FROM SALES\_DATA\_NEW WHERE PRODUCTID IN (SELECT PRODUCTID FROM TOP\_SELLERS)) S

INNER JOIN

PRODUCTS P

ON

S.PRODUCTID = P.PRODUCTID

WHERE

(YEAR (S.ORDERDATE) IN (SELECT MAX (YEAR (ORDERDATE)) FROM SALES\_DATA\_NEW ))

GROUP BY

P.PRODUCTNAME, YEAR (ORDERDATE) )

ORDER BY REVENUE DESC LIMIT 10;

**THIS ENTIRE SUBQUERY CAN  
BE SAVED AS A VIEW**

# HERE IS A PRETTY COMPLEX QUERY

```
SELECT PRODUCTNAME,REVENUE FROM  
Top_Selling_Products  
ORDER BY REVENUE DESC LIMIT 10;
```

**WE HAVE DRAMATICALLY  
REDUCED THE QUERY  
COMPLEXITY**

# YOU CAN CREATE A VIEW JUST AS YOU WOULD CREATE A TABLE FROM A SUBQUERY

```
CREATE VIEW Top_Selling_Products AS
SELECT
P.PRODUCTNAME, YEAR (ORDERDATE) , SUM (REVENUE)
FROM
(SELECT * FROM SALES_DATA_NEW WHERE PRODUCTID IN (SELECT PRODUCTID FROM
TOP_SELLERS)) S
INNER JOIN
PRODUCTS P
ON
S.PRODUCTID = P.PRODUCTID
WHERE
(YEAR(S.ORDERDATE) IN (SELECT MAX (YEAR (ORDERDATE)) FROM SALES_DATA_NEW ))
GROUP BY
P.PRODUCTNAME, YEAR (ORDERDATE) ;
```

# VIEWS

VIEWS ARE DIFFERENT FROM  
TEMPORARY TABLES

# VIEWS

WHEN A VIEW IS CREATED, IT **JUST**  
**SAVES THE QUERY** IN HIVE'S METASTORE

THE QUERY **IS NOT EXECUTED** AT THE  
TIME OF VIEW CREATION



# THE QUERY IS NOT EXECUTED AT THE TIME OF VIEW CREATION

```
hive> Create VIEW Top_Selling_Products AS
> SELECT
> p.ProductName, YEAR(orderdate), SUM(revenue)
> FROM
> (select * from Sales_Data_NEW where ProductID in (select PRODUCTID from
> top_sellers)) s
> INNER JOIN
> Products p
> ON
> s.ProductID = p.ProductID
> WHERE
> (YEAR(s.orderdate) in (SELECT max(YEAR(orderdate)) FROM Sales_Data_NEW ))
> GROUP BY
> p.ProductName, YEAR(orderdate);
```

OK

Time taken: 0.101 seconds

# VIEWS

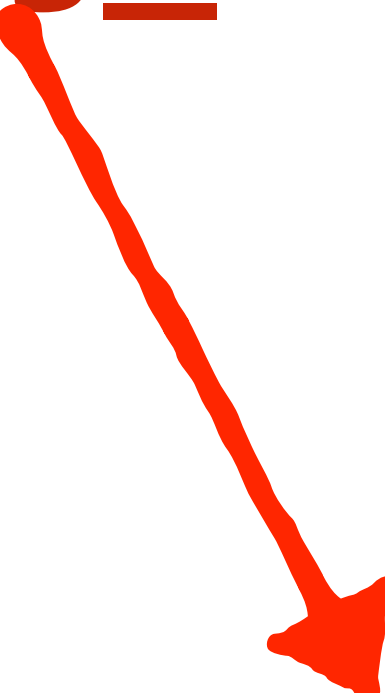
YOU CAN TREAT A VIEW LIKE A  
TABLE AND QUERY IT

```
select ProductName  
from Top_Selling_Products;
```

# VIEWS

```
select ProductName  
from Top_Selling_Products;
```

IN THE BACKGROUND, HIVE WILL  
REPLACE THE VIEW NAME  
WITH THE ACTUAL QUERY

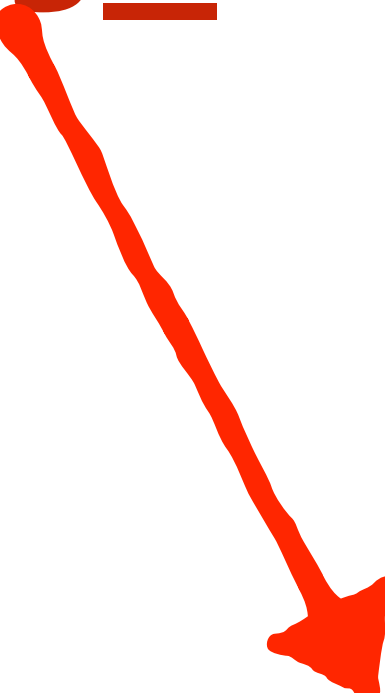


```
SELECT  
P.PRODUCTNAME, YEAR (ORDERDATE) AS  
ORDERDATE, SUM (REVENUE) AS REVENUE  
FROM  
(SELECT * FROM SALES_DATA_NEW WHERE PRODUCTID IN (SELECT PRODUCTID FROM  
TOP_SELLERS)) S  
INNER JOIN  
PRODUCTS P  
ON  
S.PRODUCTID = P.PRODUCTID  
WHERE  
(YEAR (S.ORDERDATE) IN (SELECT MAX (YEAR (ORDERDATE)) FROM  
SALES_DATA_NEW ))  
GROUP BY  
P.PRODUCTNAME, YEAR (ORDERDATE);
```

# VIEWS

```
select ProductName  
from Top_Selling_Products;
```

THE VIEW'S QUERY IS  
TREATED AS A SUBQUERY



```
SELECT  
P.PRODUCTNAME, YEAR (ORDERDATE) AS  
ORDERDATE, SUM (REVENUE) AS REVENUE  
FROM  
(SELECT * FROM SALES_DATA_NEW WHERE PRODUCTID IN (SELECT PRODUCTID FROM  
TOP_SELLERS)) S  
INNER JOIN  
PRODUCTS P  
ON  
S.PRODUCTID = P.PRODUCTID  
WHERE  
(YEAR (S.ORDERDATE) IN (SELECT MAX (YEAR (ORDERDATE)) FROM  
SALES_DATA_NEW ))  
GROUP BY  
P.PRODUCTNAME, YEAR (ORDERDATE);
```

# VIEWS

VIEWS ARE INCLUDED IN THE OUTPUT OF “SHOW TABLES” COMMAND

```
Time taken: 0.101 seconds
hive> show tables;
OK
campus_housing
customer_table
customer_table_json
emailaddresses
movies
names
products
revenue
reviews
sales_data
sales_data_date_product_partition
sales_data_new
sales_data_product_partition
sales_data_without_partition
students
top_sellers
top_selling_products
```



# ONE CAN GET EXTRA INFORMATION ABOUT THE VIEW BY USING “DESCRIBE EXTENDED”

```
hive> DESCRIBE EXTENDED top_selling_products;
OK
productname          varchar(30)
_c1                   int
_c2                   decimal(20,2)

Detailed Table Information      Table(tableName:top_selling_products, dbName:default, owner:navdeepsingh, createTime:146
4445524, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[FieldSchema(name:productname, type:varchar(30), comme
nt:null), FieldSchema(name:_c1, type:int, comment:null), FieldSchema(name:_c2, type:decimal(20,2), comment:null)], locat
ion:null, inputFormat:org.apache.hadoop.mapred.SequenceFileInputFormat, outputFormat:org.apache.hadoop.hive ql.io.HiveSe
quenceFileOutputFormat, compressed:false, numBuckets:-1, serdeInfo:SerDeInfo(name:null, serializationLib:null, parameter
s:{}), bucketCols:[], sortCols:[], parameters:{}, skewedInfo:SkewedInfo(skewedColNames:[], skewedColValues:[], skewedCol
ValueLocationMaps:{}), storedAsSubDirectories:false), partitionKeys:[], parameters:{transient_lastDdlTime=1464445524}, v
iewOriginalText:SELECT
p.ProductName, YEAR(orderdate), SUM(revenue)
FROM
(select * from Sales_Data_NEW  where ProductID in (select PRODUCTID from
top_sellers)) s
INNER JOIN
Products p
ON
s.ProductID = p.ProductID
WHERE
(YEAR(s.orderdate) in (SELECT max(YEAR(orderdate)) FROM Sales_Data_NEW ))
GROUP BY
p.ProductName, YEAR(orderdate), viewExpandedText:SELECT
`p`.`productname`, YEAR(`s`.`orderdate`), SUM(`s`.`revenue`)
FROM
(select `sales_data_new`.`productid`, `sales_data_new`.`orderdate`, `sales_data_new`.`revenue`, `sales_data_new`.`storei
d` from `default`.`Sales_Data_NEW`  where ProductID in (select `top_sellers`.`productid` from
`default`.`top_sellers`)) `s`
INNER JOIN
`default`.`Products` `p`
ON
`s`.`productid` = `p`.`productid`
WHERE
(YEAR(`s`.`orderdate`) in (SELECT max(YEAR(`sales_data_new`.`orderdate`)) FROM `default`.`Sales_Data_NEW` ))
GROUP BY
`p`.`productname`, YEAR(`s`.`orderdate`), tableType:VIRTUAL_VIEW)
```



```
Detailed Table Information      Table(tableName:top_selling_products, dbName:default, owner:navdeepsingh, cr
4445524, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[FieldSchema(name:productname, type:varcha
nt:null), FieldSchema(name:_c1, type:int, comment:null), FieldSchema(name:_c2, type:decimal(20,2), comment:n
ion:null, inputFormat:org.apache.hadoop.mapred.SequenceFileInputFormat, outputFormat:org.apache.hadoop.hive.
quenceFileOutputFormat, compressed:false, numBuckets:-1, serdeInfo:SerDeInfo(name:null, serializationLib:nul
s:{}), bucketCols:[], sortCols:[], parameters:{}, skewedInfo:SkewedInfo(skewedColNames:[], skewedColValues:[
ValueLocationMaps:{}), storedAsSubDirectories:false), partitionKeys:[], parameters:{transient lastDdlTime=14
```

```
iewOriginalText:SELECT
p.ProductName, YEAR(orderdate), SUM(revenue)
FROM
(select * from Sales_Data_NEW  where ProductID in (select PRODUCTID from
top_sellers)) s
INNER JOIN
Products p
ON
s.ProductID = p.ProductID
WHERE
(YEAR(s.orderdate) in (SELECT max(YEAR(orderdate)) FROM Sales_Data_NEW ))
GROUP BY
p.ProductName, YEAR(orderdate), viewExpandedText:SELECT
`p`.`productname`, YEAR(`s`.`orderdate`), SUM(`s`.`revenue`)
FROM
(select `sales_data_new`.`productid`, `sales_data_new`.`orderdate`, `sales_data_new`.`revenue`, `sales_data
d` from `default`.`Sales_Data_NEW`  where ProductID in (select `top_sellers`.`productid` from
`default`.`top_sellers`)) `s`
INNER JOIN
`default`.`Products` `p`
ON
`s`.`productid` = `p`.`productid`
```

**THE QUERY TEXT IS PART OF THE VIEW  
METADATA**



# VIEWS

VIEWS IN HIVE ARE "READ-ONLY", IT MAY NOT BE USED AS THE TARGET OF LOAD/INSERT



# VIEWS

METADATA OF VIEWS CAN BE ALTERED  
USING THE ALTER VIEW COMMAND

# VIEWS

METADATA OF VIEWS CAN BE **ALTERED** USING **ALTER VIEW** COMMAND

```
ALTER VIEW Top_Selling_Products AS
SELECT
P.PRODUCTNAME, YEAR (ORDERDATE) AS ORDERDATE, SUM (REVENUE) AS REVENUE
FROM
(SELECT * FROM SALES_DATA_NEW WHERE PRODUCTID IN (SELECT
PRODUCTID FROM
TOP_SELLERS)) S
INNER JOIN
PRODUCTS P
ON
S.PRODUCTID = P.PRODUCTID
GROUP BY
P.PRODUCTNAME, YEAR (ORDERDATE) ;
```

**WE HAVE CHANGED THE  
QUERY**

# VIEWS

VIEWS CAN COME IN PRETTY HANDY FOR

1. REDUCING QUERY COMPLEXITY
2. RESTRICTING ACCESS TO DATA
3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

# VIEWS

## 1. REDUCING QUERY COMPLEXITY

**WE'VE ALREADY SEEN HOW VIEWS  
HELP WITH THIS**

3. CONSTRUCTING DIFFERENT LOGICAL  
TABLES USING 1 PHYSICAL TABLE

# VIEWS

VIEWS CAN COME IN PRETTY HANDY FOR

1. REDUCING QUERY COMPLEXITY

2. RESTRICTING ACCESS TO DATA

YOU CAN CREATE A VIEW OVER A TABLE WITH

3. CONSTRUCTING DIFFERENT LOGICAL

A SUBSET OF COLUMNS

ROWS WHICH SATISFY CERTAIN CONDITIONS

# VIEWS

## 2. RESTRICTING ACCESS TO DATA

SAY WE HAD A TABLE WITH THE FOLLOWING COLUMNS

OrderID	CustomerID	ProductCategory	OrderDate	Revenue
---------	------------	-----------------	-----------	---------

A LOT OF FOLKS WANT TO  
QUERY THIS TABLE

ENGINEERS, MARKETERS, CATEGORY  
MANAGERS



# VIEWS

## 2. RESTRICTING ACCESS TO DATA

OrderID	CustomerID	ProductCategory	OrderDate	Revenue
---------	------------	-----------------	-----------	---------

LET'S CONSIDER  
CATEGORY  
MANAGERS

WE DON'T WANT THEM TO HAVE  
ACCESS TO THE CUSTOMER ID COLUMN

WE ALSO WANT TO ALLOW THE USER TO  
ONLY SEE THE REVENUE FOR THEIR CATEGORY

# VIEWS

## 2. RESTRICTING ACCESS TO DATA

OrderID	CustomerID	ProductCategory	OrderDate	Revenue
---------	------------	-----------------	-----------	---------

LET'S CONSIDER  
CATEGORY  
MANAGERS

```
CREATE VIEW  
OrderData_Books_Category AS  
SELECT ORDERID, PRODUCTCATEGORY,  
ORDERDATE, REVENUE FROM  
ORDERDATA WHERE  
PRODUCTCATEGORY="BOOKS";
```



# VIEWS

## 2. RESTRICTING ACCESS TO DATA

```
CREATE VIEW OrderData_Books_Category AS  
SELECT ORDERID, PRODUCTCATEGORY,  
ORDERDATE, REVENUE FROM  
ORDERDATA WHERE  
PRODUCTCATEGORY="BOOKS";
```

**THE VIEW DOES NOT HAVE THE  
CUSTOMERID COLUMN**

# VIEWS

## 2. RESTRICTING ACCESS TO DATA

```
CREATE VIEW OrderData_Books_Category AS  
SELECT ORDERID, PRODUCTCATEGORY,  
ORDERDATE, REVENUE FROM  
ORDERDATA WHERE  
PRODUCTCATEGORY="BOOKS";
```

THE VIEW ONLY HAS THE DATA FOR  
PRODUCT CATEGORY "BOOKS"

# VIEWS

## 2. RESTRICTING ACCESS TO DATA

```
CREATE VIEW OrderData_Books_Category AS  
SELECT ORDERID, PRODUCTCATEGORY,  
ORDERDATE, REVENUE FROM  
ORDERDATA WHERE  
PRODUCTCATEGORY="BOOKS";
```

**WE'LL GIVE THE BOOKS CATEGORY  
MANAGER ACCESS TO THIS VIEW  
INSTEAD OF THE ORDERDATA TABLE**

# VIEWS

## 2. RESTRICTING ACCESS TO DATA

WHY DIDN'T WE CREATE A SEPARATE  
TABLE INSTEAD OF A VIEW?

THE TABLE WOULD HAVE TO BE UPDATED  
WHenever THE ORIGINAL TABLE CHANGES

# VIEWS

VIEWS CAN COME IN PRETTY HANDY FOR

1. REDUCING QUERY COMPLEXITY

2. RESTRICTING ACCESS TO DATA

**3. CONSTRUCTING DIFFERENT LOGICAL  
TABLES USING 1 PHYSICAL TABLE**

# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

LET'S SAY WE HAVE THE FOLLOWING TABLE


CommunicationID	Type	Date	Attributes
-----------------	------	------	------------

THIS TABLE CONTAINS DETAILS OF  
CUSTOMER COMMUNICATIONS

# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

CommunicationID	Type	Date	Attributes
-----------------	------	------	------------




EMAIL, TWEET, CALL ETC

# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

CommunicationID	Type	Date	Attributes
-----------------	------	------	------------




**A MAP WITH DIFFERENT  
KEYS FOR DIFFERENT TYPES  
OF COMMUNICATION**



# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

CommunicationID	Type	Date	Attributes
-----------------	------	------	------------




**FOR EMAILS**

**{'EMAIL ID':, 'SUBJECT':, ...}**

# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

CommunicationID	Type	Date	Attributes
-----------------	------	------	------------



**FOR CALLS**

{'PHONE NO:', 'CALL DURATION:', ...}

# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

CommunicationID	Type	Date	Attributes
-----------------	------	------	------------

WE CAN CREATE A VIEW FOR  
EACH TYPE OF COMMUNICATION

# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

CommunicationID	Type	Date	Attributes
-----------------	------	------	------------

## FOR EMAILS

```
Create VIEW emailData AS  
SELECT Communicationid, Date,  
Attributes["emailID"] as emailID,  
Attributes["Subject"] as Subject  
from  
CommData where Type="EMAIL";
```

# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

CommunicationID	Type	Date	Attributes
-----------------	------	------	------------

## FOR EMAILS

Create VIEW emailData AS  
SELECT CommunicationID, Date, Attributes["Subject"] as Subject  
from  
CommData where Type="EMAIL";

**THE VIEW CONTAINS ONLY THE EMAIL COMMUNICATIONS**

# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

CommunicationID	Type	Date	Attributes
-----------------	------	------	------------

## FOR EMAILS

```
Create VIEW emailData AS
SELECT Communicationid, Date,
Attributes["emailID"] as emailID,
Attributes["Subject"] as Subject
from
CommData where Type="EMAIL";
```

THE VIEW HAS THE  
EMAILID AND SUBJECT  
ATTRIBUTES AS  
COLUMNS

# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

CommunicationID	Type	Date	Attributes
-----------------	------	------	------------

**FOR CALLS**

```
Create VIEW callData AS  
SELECT Communicationid, Date,  
Attributes["PhoneNo"] as PhoneNo,  
Attributes["CallDuration"] as CallDuration  
from  
CommData where Type="CALL";
```

**FOR CALLS, WE  
HAVE A DIFFERENT  
SET OF ATTRIBUTES**



# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

CommunicationID	Type	Date	Attributes
-----------------	------	------	------------

**FOR CALLS**

```
Create VIEW callData AS
SELECT Communicationid, Date,
Attributes["PhoneNo"] as PhoneNo,
Attributes["CallDuration"] as CallDuration
from
CommData where Type="CALL";
```

**THE CALLS VIEW  
HAS DIFFERENT  
COLUMNS**



# VIEWS

## 3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE

CommunicationID	Type	Date	Attributes
-----------------	------	------	------------

IN EFFECT WE CREATED DIFFERENT  
LOGICAL TABLES WHICH USE THE SAME  
UNDERLYING PHYSICAL TABLE

# VIEWS

VIEWS CAN COME IN PRETTY HANDY FOR

1. REDUCING QUERY COMPLEXITY
2. RESTRICTING ACCESS TO DATA
3. CONSTRUCTING DIFFERENT LOGICAL TABLES USING 1 PHYSICAL TABLE